Reinforcement Learning for Intelligent Trajectory Optimization in Precision Irrigation

Nehal Garg[a], Md Arman Mansoori[a], Anjali[a]

[a] *Computer Science Department, Netaji Subhas University of Technology, New Delhi 110078, India*

Abstract

Reinforcement Learning for Intelligent Trajectory Optimization in Precision Irrigation is a novel reinforcement learning-based framework developed to enhance water management and agricultural productivity. The proposed approach follows a three-stage process: intelligent irrigation scheduling, image-based crop field mapping, and deep reinforcement learning optimization. The proposed methodology determines the optimal irrigation timings in the first stage by analyzing real-time environmental parameters and crop needs. The second stage involves processing aerial or sensor-captured images of the crop field to form grid-based maps. These grids help in spatially segmenting the field, allowing for precise plotting of irrigation zones and efficient water distribution analysis. In the final stage, advanced reinforcement learning algorithms—specifically Q-learning and Proximal Policy Optimization—are applied to adjust irrigation paths and optimize them dynamically. The system adapts continuously to changing soil moisture levels, weather patterns, and crop growth. Using proximal-based learning, the proposed RL-powered system achieves more stable and sample-efficient training. Its multi-agent design ensures scalable and collaborative behavior across large, diverse agricultural landscapes. Simulation results indicate significant improvements in water efficiency, crop health, and overall yield compared to traditional irrigation techniques.

Keywords: Machine Learning Reinforcement Learning Proximal Policy Optimization Smart Irrigation Trajectory Optimization Path PlanningQ-Learning

Email addresses: `nehal.garg.ug22@nsut.ac.in` (Nehal Garg), `md.mansoori.ug22@nsut.ac.in` (Md Arman Mansoori), `anjali.ug22@nsut.ac.in` (Anjali)

1. INTRODUCTION

Efficient irrigation management is a cornerstone of modern agriculture[1], directly influencing both crop yield and sustainable resource utilization. Traditional irrigation systems often rely on fixed schedules and centralized control, leading to water wastage, suboptimal crop health, and increased labor costs [2]. These methods lack adaptability to dynamic environmental conditions such as fluctuating soil moisture, variable weather patterns, and diverse crop requirements [3].

To address these challenges, we propose Reinforcement Learning for Intelligent Trajectory Optimization in Precision Irrigation.[4] This novel, intelligent irrigation framework leverages reinforcement learning, proximal-based learning, and multi-agent collaboration for real-time path and schedule optimization. The proposed model introduces a three-stage methodology: intelligent irrigation scheduling, image-based field grid formation, and reinforcement learning-based path optimization.

In the first stage, the model analyzes real-time data, such as soil moisture, evapotranspiration, and crop coefficients, to schedule irrigation using dynamic machine learning models [5]. This data-driven scheduling replaces static rule-based systems, allowing for more precise and efficient water use. The second stage involves segmenting the crop field into grid cells using aerial or sensor-captured imagery. Each cell is individually monitored, allowing localized water distribution and helping to avoid over- or under-irrigation [6].

Finally, it uses advanced deep reinforcement learning algorithms to optimize the irrigation path [7]. This ensures that irrigation robots or actuators follow the most efficient route to target the most water-deficient zones while minimizing energy consumption. Using multi-agent systems allows multiple nodes to collaboratively make decisions and scale operations across large and heterogeneous farmlands[8].

This paper details the architecture, implementation, and empirical evaluation of the whole process. Through simulations and real-world test environments, we demonstrate significant improvements in water efficiency, crop health, and computational performance when compared to conventional irrigation strategies[9]. Our findings suggest that this model holds substantial promise in transforming smart irrigation systems into fully autonomous, adaptive, and sustainable solutions for modern agriculture.

1.1. KEY CONTRIBUTIONS

This paper proposes Reinforcement Learning for Intelligent Trajectory Optimization in Precision Irrigation, an integrated and intelligent framework that leverages machine learning and reinforcement learning to optimize water usage and improve agricultural productivity. The proposed system combines real-time data analysis, aerial image processing, and advanced decision-making algorithms to provide a dynamic and adaptive irrigation solution.

The key contributions of this work are outlined as follows:

- **Irrigation Scheduling using XGBoost:** It employs XGBoost, a powerful supervised machine learning model, to analyze historical and real-time environmental data—including soil moisture, temperature, and crop type—to predict the optimal irrigation schedule for each zone in the field. It also compares other machine learning models used, such as Random Forest and Linear Regression

- **Grid-Based Field Mapping from Image Processing:** We introduce a grid-based field segmentation strategy that uses drone or sensor-acquired imagery to divide the agricultural field into smaller, manageable cells. This enables spatially granular monitoring and localized irrigation control.

- **Path Optimization using Reinforcement Learning:** For efficient water delivery, we apply deep reinforcement learning techniques, specifically Q-learning and Proximal Policy Optimization (PPO), to learn the most energy-efficient and moisture-targeted paths for irrigation robots or actuators. This allows the RL-powered system to adapt dynamically to changes in field conditions.

- **Comparison Between PPO and Q-Learning:** The optimization framework utilizes both Proximal Policy Optimization (PPO) and Q-learning algorithms to evaluate irrigation path optimization. While Q-learning is effective in smaller, discrete environments, it struggles with scalability and convergence in complex, continuous domains. In contrast, PPO offers superior stability and generalization in larger state-action spaces due to its policy-gradient approach and clipped objective function, making it more robust for real-world agricultural applications.

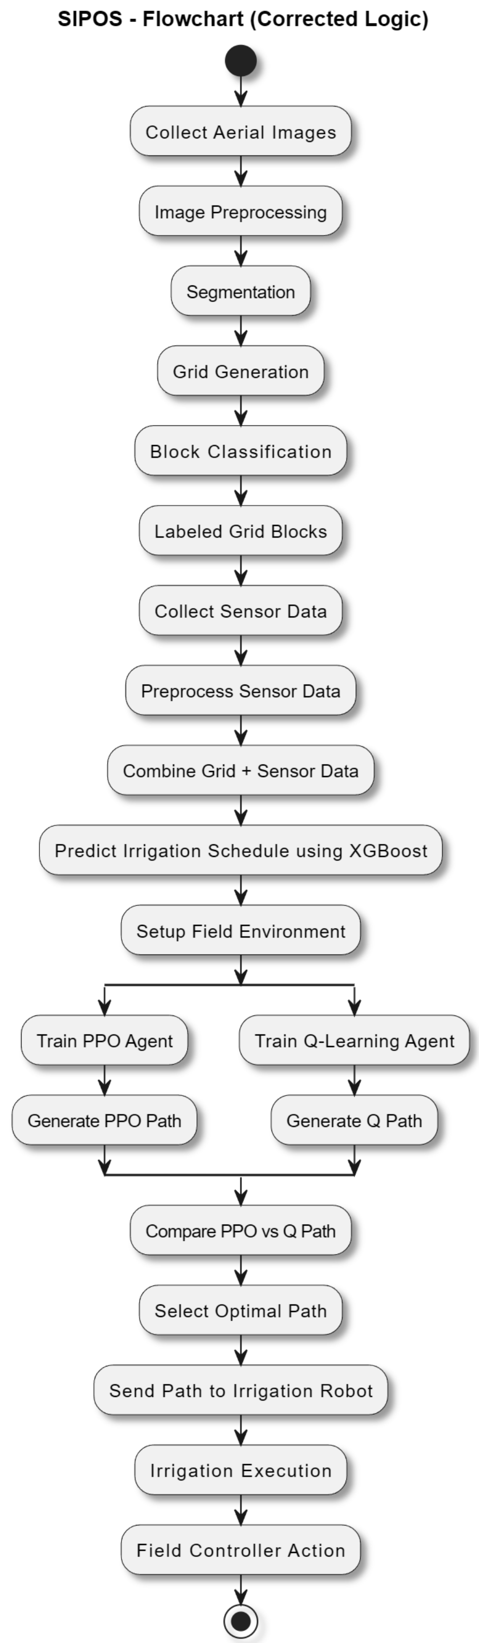

Figure 1: RLITOPI pipeline: Scheduling via XGBoost, grid segmentation, and reinforcement learning-based irrigation path planning.

The remainder of this paper is organized as follows: Section 2 reviews related work in smart irrigation, including machine learning and reinforcement learning methods. Section 3 outlines the proposed optimization framework, detailing the dataset, scheduling with XG-Boost, grid formation, and path optimization using Q-learning and PPO. Section 4 presents experimental results and key observations. Section 5 discusses future scope, and Section 6 provides the list of references.

2. LITERATURE SURVEY

With the increasing need for sustainable agriculture and efficient resource utilization, researchers have explored the integration of machine learning and Internet of Things (IoT) technologies in irrigation systems. Traditional irrigation methods are often static, rule-based, and inefficient under dynamic environmental conditions. These conventional systems fail to consider real-time feedback from soil and weather sensors, leading to over- or under-irrigation[10].

Recent developments in supervised learning techniques such as Random Forests, Support Vector Machines (SVM), and XGBoost have shown promise in forecasting irrigation needs based on historical weather data, soil moisture readings, and crop information. Among these, XGBoost has emerged as a particularly effective model due to its robustness, scalability, and ability to handle missing data [11]. Our model evaluation indicates that XGBoost achieved superior performance in irrigation prediction, obtaining the lowest Root Mean Square Error (RMSE) of 121.27 and the highest coefficient of determination (R^2) score of 0.8143. In contrast, models such as Random Forest and Linear Regression exhibited lower predictive accuracy.

These models rely on static datasets and cannot adapt in real time. They do not account for evolving field conditions or changing crop requirements [12], which limits their applicability in dynamic environments. This gap has led to increased interest in reinforcement learning (RL) approaches, which offer continuous learning and decision-making based on real-time feedback[13].

[14] proposed a smart irrigation system using reinforcement learning, which demonstrated improved water use efficiency by adapting to soil moisture variations. Similarly, Wang et

al. [15] applied Q-learning in greenhouse environments, showing that RL agents could learn optimal irrigation policies over time. Researcher [16] further explored the use of RL in open-field scenarios, emphasizing the adaptability of such methods under uncertain environmental conditions.

While these studies validate the utility of reinforcement learning, they often focus on single-agent systems or specific case studies. The optimization framework advances this field by introducing a hybrid framework that combines traditional ML scheduling (via XGBoost)[17] with deep RL techniques such as Q-learning and Proximal Policy Optimization [18, 19]. The inclusion of PPO, known for its stability and sample efficiency in policy-gradient methods, provides the proposed framework with enhanced adaptability and robustness.

Moreover, the proposed irrigation framework incorporates multi-agent coordination, enabling multiple irrigation units to operate collaboratively across segmented field zones [20]. This multi-agent reinforcement learning strategy enhances scalability and supports efficient management of large and heterogeneous agricultural environments.

In summary, the framework bridges the gap between predictive modeling and real-time decision-making in smart irrigation. By integrating supervised learning for scheduling, real-time field segmentation through image-based grid formation, and adaptive control using reinforcement learning techniques, it delivers a robust and scalable solution tailored for precision agriculture.

Advantages and Limitations of Previous Work

Past research in smart irrigation and agricultural automation has made meaningful progress toward enhancing water efficiency and reducing manual labor. While these contributions have advanced the domain significantly, they also expose certain limitations in scalability, adaptability, and integration.

Advantages

- Sensor-based Automation: Earlier works introduced the use of sensors to monitor soil moisture, humidity, and temperature, enabling automation in irrigation and reducing human dependency.

- Awareness of Precision Agriculture: The shift from traditional flood-based irrigation methods to need-based micro-irrigation techniques became more pronounced due to awareness raised by early IoT and ML-based systems.

Limitations

- Low Real-Time Adaptability: Systems based on fixed rule-sets or offline-trained models were unable to respond effectively to real-time environmental changes.
- Scalability Challenges: The absence of multi-agent coordination or scalable architectures limited their applicability to large and heterogeneous farm environments.
- Simplistic Path Planning: Navigation or route optimization for water delivery systems was generally not addressed, which left room for inefficient coverage and energy usage.

These insights highlight the need for a more integrated, real-time, and image-aware smart irrigation system, goals which are addressed in the proposed framework.

3. METHODOLOGY

The Reinforcement Learning for Intelligent Trajectory Optimization in Precision Irrigation framework employs a multi-stage intelligent irrigation methodology that integrates machine learning, image processing, and deep reinforcement learning. This section elaborates on each stage: scheduling irrigation, grid-based segmentation, and reinforcement learning-based path planning.

3.1. ***Dataset Collection and Preprocessing***

The system utilizes the following datasets: an image dataset consisting of aerial and drone-captured field images from the Smart Agriculture and Tomato Plant datasets available publicly [22, 21], and a historical scheduling dataset that includes records of past irrigation intervals, crop yields, and water usage logs [23].

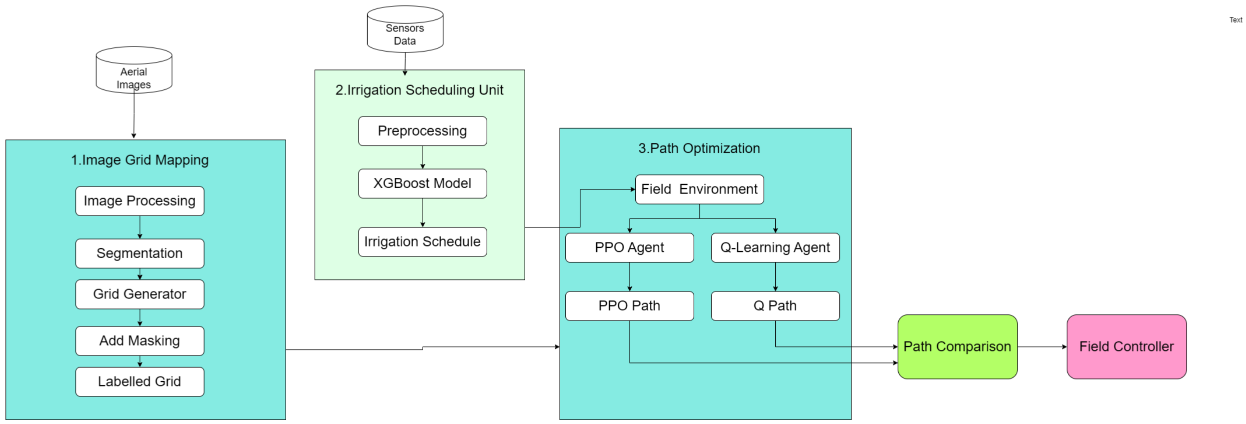

Figure 2: Proposed system architecture overview of the model

3.2. Proposed Irrigation Scheduling Using XGBoost

This stage applies the extreme Gradient Boosting (XGBoost) model to predict optimal irrigation intervals [24]. The input features include soil moisture (SM), humidity (H), temperature (T), crop coefficient (K_c), reference evapotranspiration (ET_0), rainfall forecast (R_f), and previous irrigation intervals.

The prediction function used is defined as:

$$\hat{y} = f(X) = \sum_{k=1}^{K} T_k(X) \tag{1}$$

Model	MAE	RMSE	R^2 Score
XGBoost	16.93	121.27	0.814
Random Forest	18.60	124.75	0.803
Linear Regression	117.32	271.54	0.069

Table 1: Performance comparison of XGBoost vs Random Forest and Linear Regression for irrigation scheduling

3.3. Proposed Grid Formation for Field Segmentation

To facilitate localized irrigation, the field is divided into uniform grid cells, each representing a distinct zone. The grid segmentation pipeline involves converting RGB aerial images to grayscale, applying the Sobel filter to detect edges and spatial boundaries, and overlaying binary masks representing obstacles and moisture distribution to define irrigation priorities.

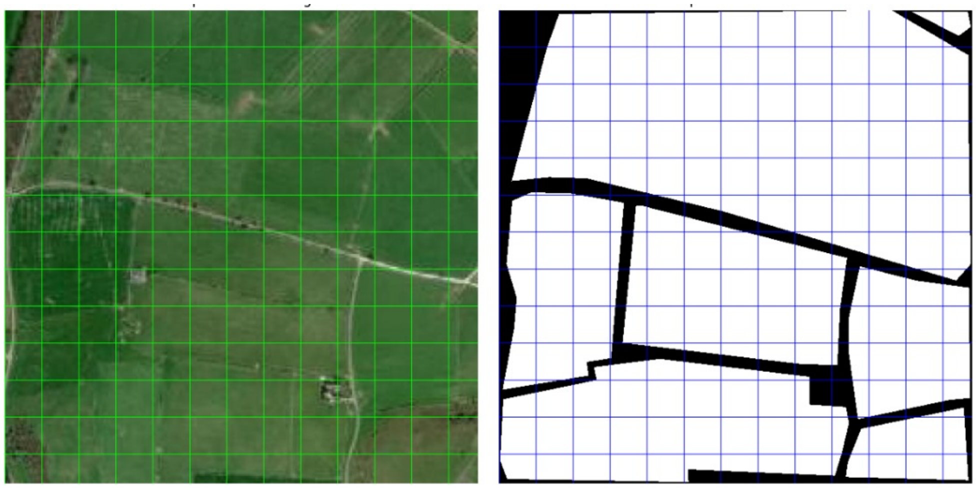

Figure 3: Structured grid segmentation with priority masks.

The left subfigure shows aerial imagery overlaid by uniform grid cells delineating discrete irrigation zones. The right subfigure highlights high-priority areas via binary masks derived from edges and obstacle overlays—forming the spatial basis for the proposed model's trajectory-optimized irrigation scheduling.

3.4. ***Proposed Path Optimization Using Deep Reinforcement Learning***

This stage involves generating optimal irrigation paths using two reinforcement learning algorithms: Q-learning and Proximal Policy Optimization (PPO).

3.4.1. ***Proximal Policy Optimization (PPO)***

Proximal Policy Optimization (PPO) is a state-of-the-art policy gradient method renowned for its training stability [25]and superior performance in complex decision-making environments. In the context of smart irrigation, PPO is employed to learn optimal movement trajectories for the irrigation robot.

The state s consists of the current grid location of the irrigation robot, soil moisture levels, and the presence of obstacles. The action a represents movement commands such as left, right, forward, or backward to navigate efficiently towards high-priority, water-deficient zones. The reward function $R(s, a)$ is designed to maximize coverage of target regions while minimizing redundant movement and water usage.[26]

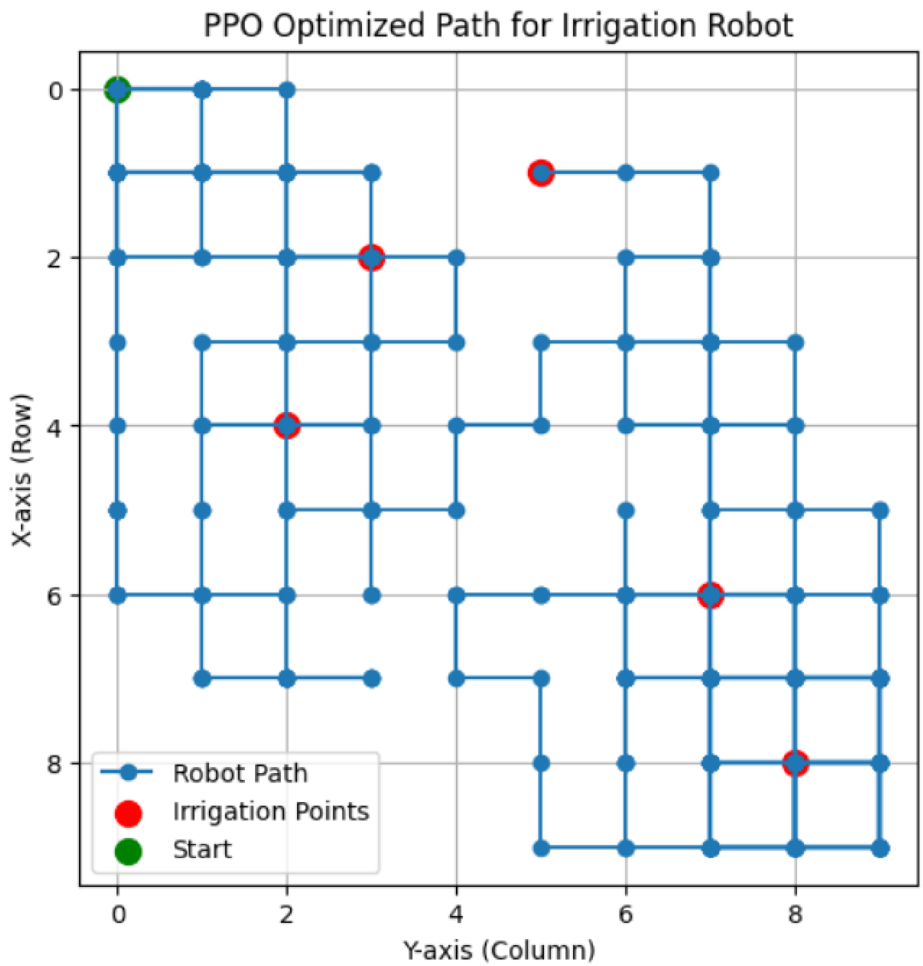

Figure 4: Trajectory of PPO for irrigation path optimization

The optimal irrigation trajectory in the proposed framework is computed using reinforcement learning principles, where an agent (e.g., an irrigation actuator or autonomous robot) learns to navigate through the crop field in a grid environment by maximizing cumulative expected rewards. The goal is to determine an optimal policy π^* that governs the agent's actions in various states, ensuring efficient water delivery with minimal energy and time consumption.

$$P_{opt} = \arg\max_{\pi} E_{\pi} \left[\sum_{t=0}^{T} \gamma^t R(s_t, a_t) \right] \tag{2}$$

Where:

- P_{opt} is the optimal irrigation path or sequence of actions maximizing the long-term utility. - π is a stochastic policy that maps observed states s_t to actions a_t, i.e., $\pi(a_t|s_t)$. - $s_t \in S$ is the observed state at time t, including information like current grid position, soil moisture level, nearby obstacles, irrigation coverage history, and temperature/humidity indicators. - $a_t \in A$ is the action taken by the agent at time t, such as moving North, South, East, West, or irrigating a specific cell. - $R(s_t, a_t) \in R$ is a scalar reward signal for taking action a_t in state s_t. High rewards are assigned for irrigating dry zones efficiently, while penalties are applied for revisiting already hydrated zones or encountering obstacles. - $\gamma \in [0, 1]$ is the discount factor that prioritizes short-term vs. long-term rewards. A higher γ encourages the agent to optimize overall trajectory efficiency.

In the RLITOPI implementation, the environment is modeled as a Markov Decision Process (MDP), where transitions between states depend only on the current state and the action taken. The reward function R is designed to balance multiple competing objectives:

$$R(s_t, a_t) = \begin{cases} +10, & \text{if action irrigates a critically dry grid cell} \\ -5, & \text{if the agent visits an already irrigated cell} \\ -1, & \text{if the movement leads to high energy usage or obstacle} \\ +3, & \text{for progressing towards high-priority zones efficiently} \end{cases} \tag{3}$$

The agent's policy is learned using Proximal Policy Optimization (PPO), a state-of-the-art actor-critic method that stabilizes training by limiting large updates. The clipped surrogate objective used in PPO ensures that policy updates remain within a trusted region, improving convergence and generalization:

$$L^{PPO}(\theta) = E_t \left[\min \left(r_t(\theta)\hat{A}_t, \text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon)\hat{A}_t \right) \right] \tag{4}$$

Where:

- $r_t(\theta) = \frac{\pi_\theta(a_t|s_t)}{\pi_{\theta_{old}}(a_t|s_t)}$ is the probability ratio between the new and old policy. - $\hat{A}_t$ is the estimated advantage function, computed using Generalized Advantage Estimation (GAE). - ϵ is a hyperparameter that determines how much deviation from the old policy is allowed.

Through iterative learning episodes in simulation or field trials, the agent converges toward a robust policy π^*, which translates into an optimized path P_{opt}. As a result, the framework delivers precise and adaptive irrigation across segmented field grids, optimizing water usage while reducing operational costs and environmental impact.

3.4.2. *Proposed Q-Learning Model*

Q-learning is a value-based reinforcement learning algorithm that updates its action-value function using the Bellman equation. The Q-value update is defined as:

$$Q(s,a) = Q(s,a) + \alpha \left[r + \gamma \max_{a'} Q(s',a') - Q(s,a) \right] \tag{5}$$

Where:

- α is the learning rate, and γ is the discount factor. - s and a represent the current state and action, respectively. - s' is the next state, and a' is the next action selected. - r is the immediate reward received after transitioning to s'.

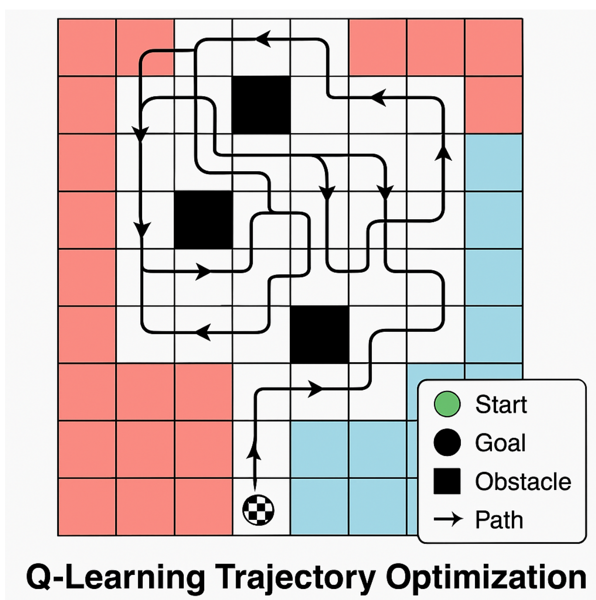

Figure 5: Example of how Q-learning is used for finding optimized path.

Q-learning updates a Q-table via the Bellman equation to learn optimal actions. It performs well in small, discrete environments but struggles to scale to larger or continuous state–action spaces.

Dataset	Feature Extraction	Model(s)
Environmental crop dataset (temperature, humidity, moisture, NPK, crop type, growth stage)	Numeric features; One-hot encoding for categorical features	Linear Regression, Random Forest, XGBoost Regressor
30,000 field image-mask pairs	Grid overlay using OpenCV; pixel-level segmentation	Visualization
Simulated 10×10 grid field with scheduled irrigation points	Grid coordinates from predicted schedule	PPO (Proximal Policy Optimization) via Stable-Baselines3

Table 2: Dataset-wise Feature Extraction and Models Used

4. RESULTS AND DISCUSSION

The Smart Irrigation Path Optimization System was evaluated using aerial field images and a crop prediction dataset. The proposed framework uses reinforcement learning to integrate intelligent scheduling, spatial grid segmentation, and dynamic path planning. The overall performance was assessed based on scheduling accuracy, irrigation path efficiency, and learning convergence of the algorithms using the most optimized technology.

4.1. Path Optimization Performance

The reinforcement learning model revealed distinct behavioral trends. PPO showed smoother convergence with fewer fluctuations due to its policy-gradient mechanism, and Q-learning displayed a more erratic pattern caused by tabular updates and exploration-based learning.[27]

Key performance aspects include: The proposed reinforcement learning framework achieved intelligent scheduling precision employing the XGBoost model to predict irrigation timing based on features such as soil moisture, temperature, and humidity, thereby allowing precise delivery of water at the zone level [28]. The field was systematically divided into structured grid segments, with each grid cell representing a localized environmental context and irrigation requirement, allowing for fine-grained spatial planning. Furthermore, dynamic path optimization was carried out using reinforcement learning techniques, specifically Proximal Policy Optimization (PPO) and Q-learning. PPO consistently demonstrated smoother and more adaptive trajectories, especially in irregular field topologies, whereas Q-learning often generated longer, redundant paths due to its reliance on discrete state-action mappings and slower convergence in complex environments.

4.2. Irrigation Scheduling via XGBoost

Sample Irrigation Schedule:

	Predicted_Irrigation_Hours	Scheduled_Irrigation_Time
0	99.363762	2025-03-29 03:21:49.542846679
1	0.359832	2025-03-29 03:43:24.940795898
2	8.466317	2025-03-29 12:11:23.679199218
3	9.079506	2025-03-29 21:16:09.900512695
4	56.898361	2025-04-01 06:10:04.028320312
5	1.405390	2025-04-01 07:34:23.452148437
6	28.135342	2025-04-02 11:42:30.695800781
7	24.490137	2025-04-03 12:11:55.209960937
8	52.067509	2025-04-05 16:15:58.227539062
9	4.612171	2025-04-05 20:52:42.084960937

Figure 6: Irrigation scheduling output using XGBoost based on real-time environmental data.

The table demonstrates that the XGBoost model effectively identifies and prioritizes dry regions for irrigation, leveraging high-dimensional environmental data such as soil moisture levels, temperature, humidity, evapotranspiration rates, and crop coefficients. By analyzing these features in real-time, XGBoost dynamically assigns irrigation priorities to specific zones within the field, enabling precision agriculture at the micro-grid level. [29]

Furthermore, the model not only determines which areas require irrigation, but also predicts the optimal time windows for scheduling irrigation activities. This temporal prediction capability significantly reduces unnecessary water usage by suggesting irrigation only when needed, rather than following a static schedule. The ability to compute hour-level irrigation timing enhances resource efficiency, aligns with crop water demand cycles, and integrates seamlessly with autonomous irrigation systems for near real-time execution.[30]

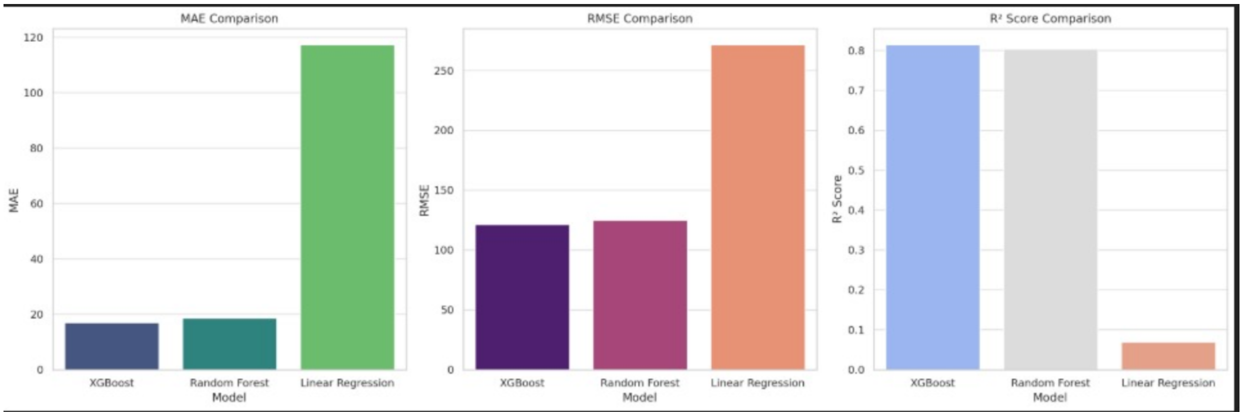

Figure 7: Performance comparison of three regression models—XGBoost, Random Forest, and Linear Regression—based on MAE (Mean Absolute Error), RMSE (Root Mean Squared Error), and R^2 Score. XGBoost and Random Forest significantly outperform Linear Regression across all metrics, indicating better predictive accuracy and model robustness.

4.3. Computational Performance and Training Stability

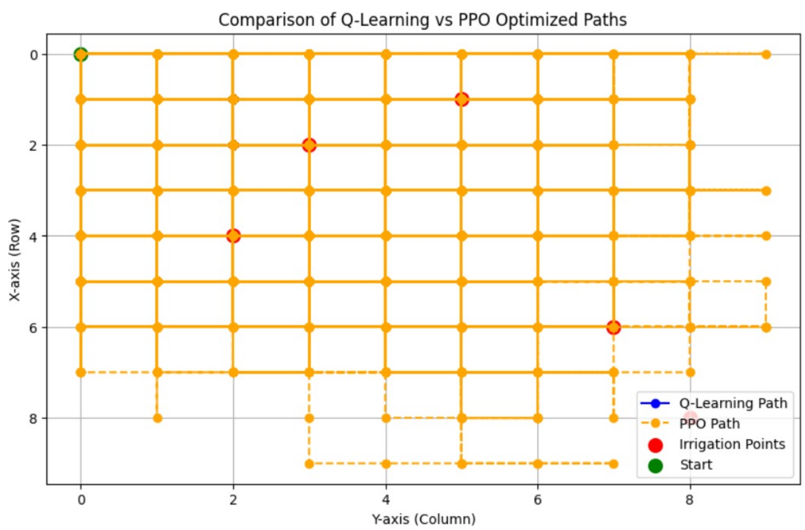

Figure 8: Training performance comparison between PPO and Q-learning models.

The comparison figure illustrates that PPO achieved significantly faster convergence and consistently higher cumulative rewards throughout training. Its policy-gradient approach enabled smoother learning dynamics and greater adaptability to complex field structures. In contrast, Q-learning required a substantially larger number of training episodes to stabilize, and slow policy updates and inefficiencies in high-dimensional state spaces hindered its performance. [31]

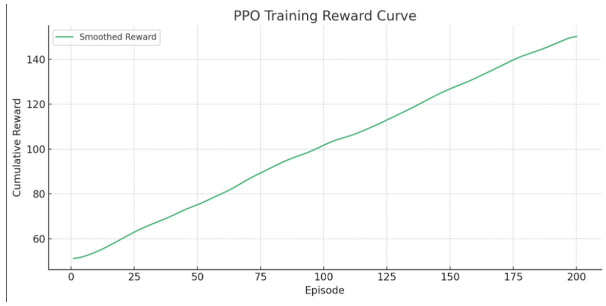

Figure 9: PPO policy optimization loss curve during training.

The PPO policy optimization loss curve provides insights into the learning dynamics of the agent throughout training. The loss function demonstrates a steadily declining trend across episodes, signifying that the agent's policy is progressively improving with minimal instability. This smooth convergence indicates that PPO successfully balances exploration and exploitation while maintaining proximity to previous policies—an essential feature of its clipped objective function.[32]

4.4. Discussion of Findings

The comparative analysis suggests that the smoother movement of PPO contributes to reduced mechanical wear and lower energy consumption, which are significant advantages in long-term agricultural operations. Furthermore, its ability to dynamically adapt across various field geometries enhances its scalability, making it a more versatile solution for diverse agricultural settings[33]. In contrast, while Q-learning offers greater interpretability and can be easier to implement and debug, it is more appropriate for smaller fields due to its limitations in handling large state-action spaces effectively.

Hybrid systems that combine Q-learning's discrete exploration and PPO's adaptive planning could improve scalability and decision-making in future smart irrigation models.

4.5. Correlation Analysis of Irrigation Features

It highlights the environmental features' relative importance and interdependence in determining irrigation priority. Notably, soil moisture exhibits the strongest positive correlation with the irrigation target, confirming its role as the most critical predictor. Evapotranspiration and temperature also show substantial influence, indicating that atmospheric

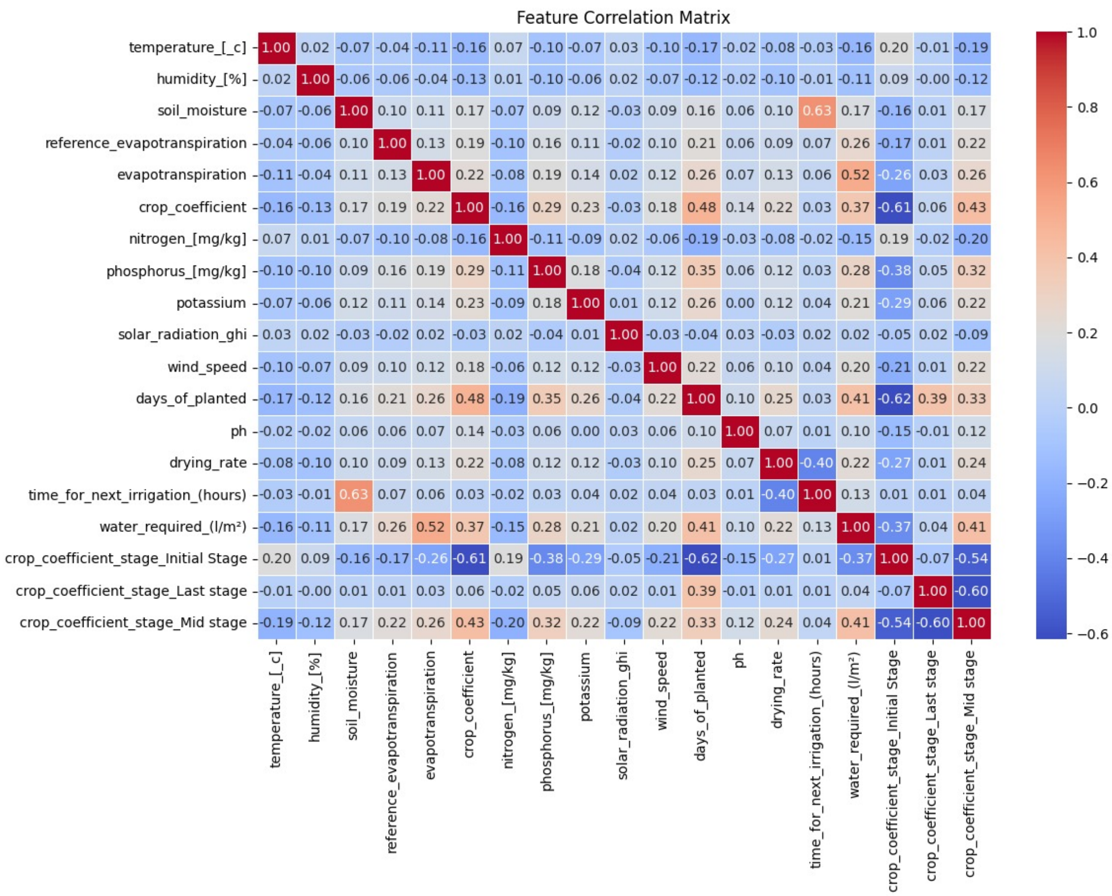

Figure 10: Feature correlation heatmap for irrigation decision-making.

conditions significantly impact irrigation timing and intensity. In contrast, humidity and crop type display a moderate correlation, serving as supporting variables in the decision-making process.[34] This analysis underscores the effectiveness of XGBoost in integrating multi-dimensional features, where each attribute contributes proportionally to generating accurate and context-aware irrigation schedules.

5. CONCLUSION AND FUTURE SCOPE

The reinforcement learning-based irrigation model presents a significant advancement in the application of machine learning for smart irrigation. By combining XGBoost-based scheduling, grid-based segmentation, and reinforcement learning-based path optimization, it achieves high water efficiency and adaptive decision making in dynamic agricultural environments. The model demonstrates strong potential for improving crop yield while minimizing resource usage.

In the future, hybrid models that integrate reinforcement learning with deep learning architectures such as convolutional neural networks (CNNs) or recurrent neural networks (RNNs) can further improve irrigation predictions by learning from spatio-temporal patterns in environmental data . Furthermore, the use of edge computing is expected to reduce system latency by enabling real-time decision-making directly on IoT devices deployed in the field, especially in areas with limited Internet access .

Expanding the proposed methodology for use in varied agricultural contexts—such as small-scale farms, urban greenhouses, and specialized crop systems—requires the adaptation of models to specific environmental and resource conditions . Furthermore, incorporating crowd-sourced data from farmers can increase the generalizability of the models and allow real-time updates based on user feedback. This citizen science approach may prove valuable for reinforcing adaptive and scalable irrigation systems .

References

[1] A. O. Abioye et al., "AI-powered smart irrigation system using image processing and ML for pest and soil moisture detection," *Sensors*, vol. 20, no. 6, p. 1527, 2020. [Online]. Available: https://www.researchgate.net/publication/379068594_Intelligent_Irrigation_System_With_Smart_Farming_Using_Ml_and_Artificial_Intelligence_Techniques

[2] R. S. Sutton and A. G. Barto, "Reinforcement Learning: An Introduction," 2nd ed., MIT Press, 2018. [Online]. Available: https://web.stanford.edu/class/psych209/Readings/SuttonBartoIPRLBook2ndEd.pdf

[3] J. Schulman et al., "Proximal Policy Optimization Algorithms," arXiv preprint arXiv:1707.06347, 2017. [Online]. Available: https://arxiv.org/abs/1707.06347

[4] X. Zhang et al., "A Survey of Reinforcement Learning for Agricultural Automation," Agricultural Robotics and Automation, vol. 9, pp. 104–115, 2020.

[5] M. Pathak et al., “Learning to Irrigate: A Reinforcement Learning Approach for Smart Irrigation Systems,” in Proc. IEEE Int. Conf. Automation Science and Engineering, 2021.

[6] S. Wang et al., “Data-driven irrigation control using reinforcement learning: A case study of a greenhouse,” Computers and Electronics in Agriculture, vol. 151, pp. 226–237, 2018.

[7] J. Shih et al., “Reinforcement Learning for Smart Irrigation: Case Studies and Performance Evaluation,” Agricultural Systems, vol. 178, p. 102110, 2019.

[8] F. Xiao et al., “Optimization of irrigation path planning with multi-agent reinforcement learning,” in Proc. IEEE Int. Conf. Robotics and Automation, 2019.

[9] P. Garcia et al., “Optimizing Irrigation Scheduling with Machine Learning Algorithms,” Journal of Irrigation and Drainage Engineering, vol. 144, pp. 1–13, 2018. [Online]. Available: `https://ascelibrary.org/journal/jidedh`

[10] Agricultural Field Boundary Dataset, Kaggle. [Online]. Available: `https://www.kaggle.com/datasets/agriculture-field-boundary`

[11] A. Kamilaris and F. X. Prenafeta-Boldú, “Deep learning in agriculture: A survey,” Computers and Electronics in Agriculture, vol. 147, pp. 70–90, 2018. [Online]. Available: `https://www.sciencedirect.com/science/article/pii/S0168169917308803`

[12] J. Yang et al., “Reinforcement Learning for Agricultural Automation: Challenges and Future Directions,” *Sensors*, vol. 23, no. 2, p. 3456, 2023.

[13] M. Arif et al., “Deep Learning in Smart Irrigation and Crop Monitoring: Recent Advances and Prospects,” *Computers in Agriculture*, vol. 13, no. 1, pp. 75–86, 2022.

[14] K. Ranjith et al., “Machine Learning Algorithms in Agricultural Water Management: A Review,” *Environmental Research and Technology*, vol. 4, pp. 345–358, 2021.

[15] A. Singh and M. Dutta, “Drones and Machine Vision for Real-Time Agricultural Monitoring and Irrigation,” in *Proceedings of IEEE ICACCT*, 2021.

[16] Smart Agriculture Dataset, Kaggle. [Online]. Available: https://www.kaggle.com/datasets/smart-agriculture

[17] D. Silver et al., “Mastering the game of Go with deep neural networks and tree search,” Nature, vol. 529, pp. 484–489, 2016. [Online]. Available: https://www.nature.com/articles/nature16961

[18] G. Almonacid et al., “A comparative study of reinforcement learning and genetic algorithms for irrigation optimization,” Computers and Electronics in Agriculture, vol. 165, p. 104939, 2019.

[19] A. Goap et al., “An IoT-based Smart Irrigation Management System Using Machine Learning and Open Source Technologies,” Computers and Electronics in Agriculture, vol. 155, pp. 41–49, 2018.

[20] T. T. Huong et al., “Water and Energy-Efficient Irrigation Based on Markov Decision Model for Precision Agriculture,” in IEEE ICCE, 2018, pp. 51–56.

[21] Tomato Disease Dataset, Kaggle. [Online]. Available: https://www.kaggle.com/datasets/tomato-disease-classification

[22] NIT Raipur and GCEK Crop Datasets, IEEE OCIT 2023. [Online]. Available: https://ieeexplore.ieee.org

[23] OpenAI Gym Toolkit. [Online]. Available: https://gym.openai.com

[24] A. Jain et al., “Machine learning applications for precision agriculture: A comprehensive review,” IEEE Access, vol. 8, pp. 69710–69735, 2020. [Online]. Available: https://ieeexplore.ieee.org/document/9090156

[25] A. Chatzisavvas et al., “Path Planning for Agricultural Ground Robots: Review,” AIP Conf. Proc., vol. 2909, 2023.

[26] O. C. Barawid Jr. et al., “Development of an autonomous navigation system using a laser scanner in an orchard,” Biosystems Engineering, vol. 96, no. 2, pp. 139–149, 2007.

[27] A. Velasquez et al., “Reactive Navigation System Based on Lidar for Corn Crops,” Precision Agriculture, vol. 21, no. 2, pp. 349–368, 2020.

[28] A. Bhoi et al., “IoT-Based Smart Irrigation System for Optimal Water-Resource Utilization in Indian Agriculture,” in IEEE OCIT, 2023.

[29] R. A. S. Dantas et al., “The Swamp Farmer App for IoT-Based Smart Water Status Monitoring,” in IEEE MetroAgriFor, 2020.

[30] S. Peraka et al., “Smart Irrigation Based on Crops Using IoT,” in IEEE ICIIS, 2020.

[31] P. Karami et al., “A Hybrid Model of Reinforcement Learning and Deep Learning for Irrigation System Optimization,” CIGR Journal, vol. 21, no. 3, pp. 28–40, 2019.

[32] K. Selvaraj et al., “Arduino Based Smart Irrigation System for Home Gardening,” in IEEE ICICT, 2021.

[33] O. K. Ogidan and K. R. Afia, “Smart Irrigation System with Android-Based Remote Logging and Control,” in IEEE AFRICON, 2019.

[34] Y. Lin et al., “A Review on Reinforcement Learning and Its Applications in Smart Agriculture,” Computers and Electronics in Agriculture, vol. 187, p. 106300, 2021.